

## Lec 14:- handling node failures in replication protocols.

When do you need consensus?

8:10

→ You have a bunch of processes, and...

(i) need to make sure they deliver the same messages in the same order.

(you need them all to keep track of which processes exist. You would need them all to keep a list of which processes exist.

And whenever anyone dies or leaves for whatever reason, everyone's list has to get updated)

It turns out that there are all classic distributed system problems.

and they have names

PTC

(ii) you need them all to know which processes exist and which are running.

(iii) You want one process to play a distinguished role (and you want other processes to know about it).

(iv) You want them to take turns getting mutually exclusive access to some resource.

(v) They are all participating in a transaction and you need them to agree on whether it's committed/aborted.

(i) → Totally Ordered Broadcast problem  
(also called atomic broadcast)

(ii) → Group membership problem  
(or failure detector)

(iii) → Leader ~~detection~~ election problem.

(iv) → Distributed mutual exclusion problem

(v) → Distributed Transaction commit problem.

So, what do all these problems have in common?

→ They all involve this bunch of processes trying to co-ordinate with one another and reach agreement about something.

The thing they are trying to agree on can vary

But all of these are going to boil down to the consensus problem,

(which is getting a group of processes to co-ordinate.

and agree on something.

Note - The thing that makes consensus hard in  
DL is faults (crash faults & omission  
faults which we are going to consider)


## Consensus

→ So, we have some processes and we want  
to get them agree.

And to make things simple, let's say that  
the thing that they are trying to agree on  
is just one bit

(zero or one.)

So, every process gets to have a say,  
and every process gets to contribute  
one bit → zero or one






The role of the consensus is to make all  
processes agree on a single ~~to~~ ~~the~~

either

or



and indeed this notion of majority is  
~~going~~ is important

Note- You have to keep in mind is there isn't  
some external process that looks at these  
numbers and makes the choice.

Rather, it's ~~there~~ three processes  
coming in without any knowledge  
of what number the other ones  
picked that have to coordinate  
among themselves and reach a point  
where they all agree and they have  
to do it in the presence of  
failure any asynchrony.

Properties that algorithms try to satisfy:-

(i) Termination:-

Each correct process eventually decides a value.

(ii) Agreement:-

all correct processes agree on the same value.

(iii) Validity or Integrity or nontriviality  
→ the agreed upon values must be one of the proposed values.

We said that consensus algorithm try to satisfy these properties.

But it turns out that no consensus algorithm actually does satisfy them all, especially in an asynchronous network model.

## Paxos

— It is a consensus algorithm invented by "Leslie Lamport"

→ Paxos is a family of algorithms.

— for now, we are just going to talk about

### "Vanilla Paxos"

→ There are three roles that a Paxos can play.

- proposer — propose values

- acceptor — contribute to a decision, from among the proposed values.

- learner — learn the agreed upon value.

Note: In practice, a single process might can take multiple roles or even all of them.

but when you are learning Paxos it is good to assume that one process can take one role at a time.

• One thing that every node is to know that what is the "majority" of acceptor is.

(eg if 3 acceptors then majority is 2.  
if 5 " " " " " 3.

So, all nodes need to know this in advance.

For now, the protocol we are looking at is for deciding on a single value.

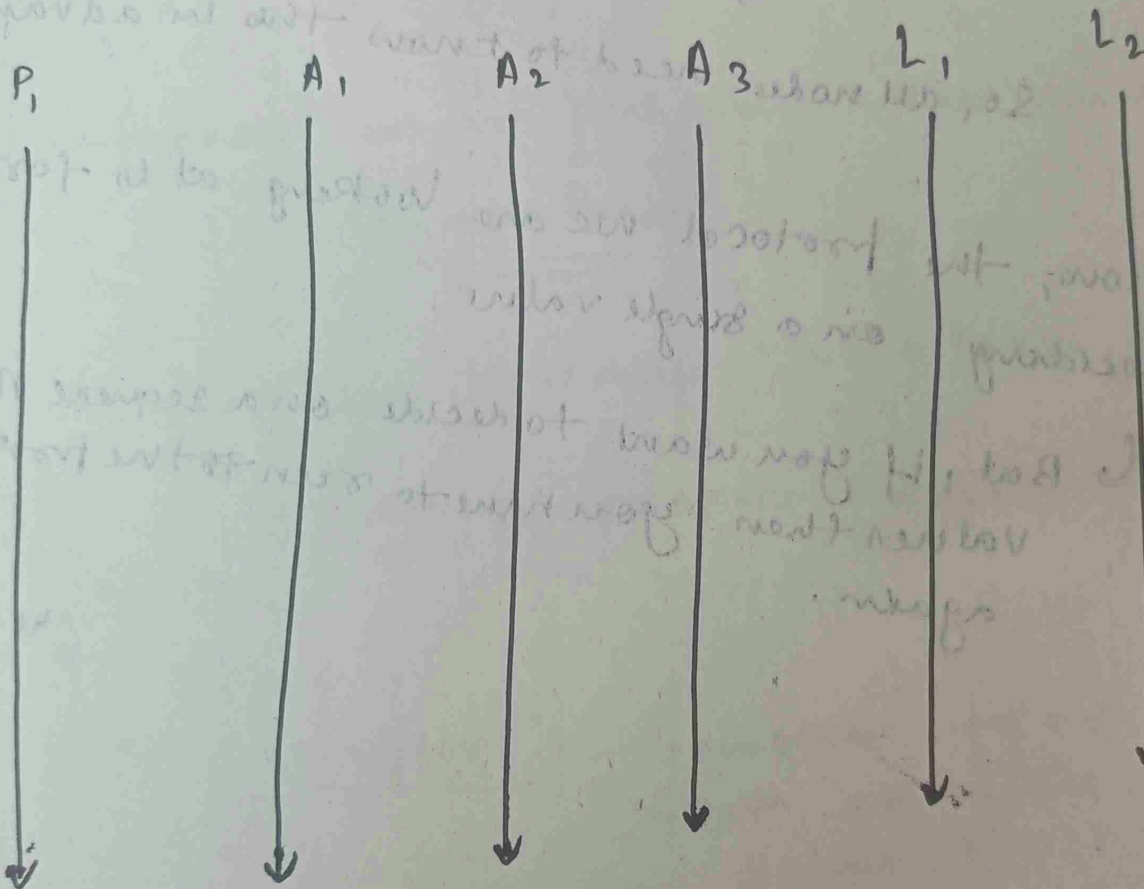
↳ But, if you want to decide on a sequence of values then you have to run the protocol again.

40:00



let's look at an actual run of the algo:

- let's say there is one proposer  $\rightarrow P_1$
- 3 acceptors ✓
- 2 learners.
- majority  $\rightarrow 2$



let's say our proposer  $P_1$  wants to propose a value.

So, the first thing it had to do is to send what's called a "Prepare message" to atleast a majority of acceptors.



So,  $P_i$  is going to send a msg "prepare", and that msg is going to come along with "proposed number".

Notes: The "proposal number" need to be unique

(And if you had multiple proposers, you might establish rules in advance,

like eg one proposer can only use odd numbered proposal nos.

and other proposer can only use even numbered proposal number)

))

That has to be established in advance

→ And the proposal number has to be higher than any of the proposal number that proposer has tried before.

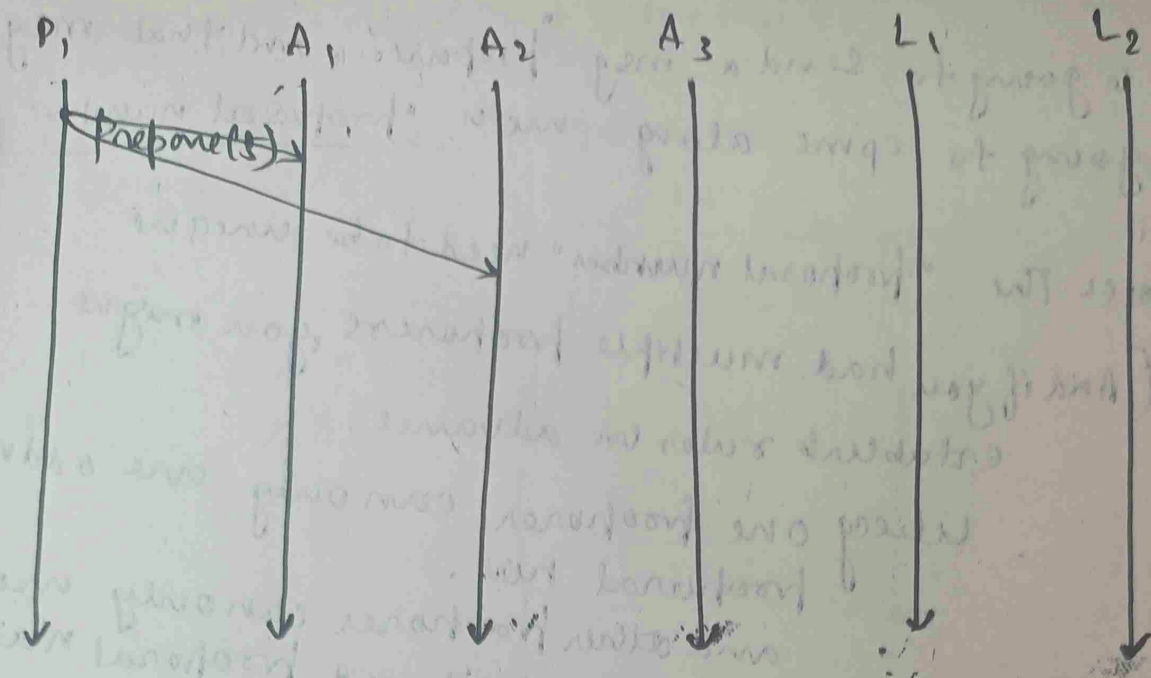
(So, it has to be able to store the proposal numbers it's used and make sure not to reuse them)

So, proposer 1 ( $P_1$ ) sends a prepare message with a proposal number.

(Let's say the proposal number is 5)

and  $P_1$  is going to send it to majority of acceptors.

(Let's say it goes to  $A_1$  &  $A_2$ )



→ when an acceptor gets a prepare message,  
 it looks at that proposal number, and it does this  
 check! -

"Did I previously promise to ignore  
 requests with this proposal number?"

If it did      If no

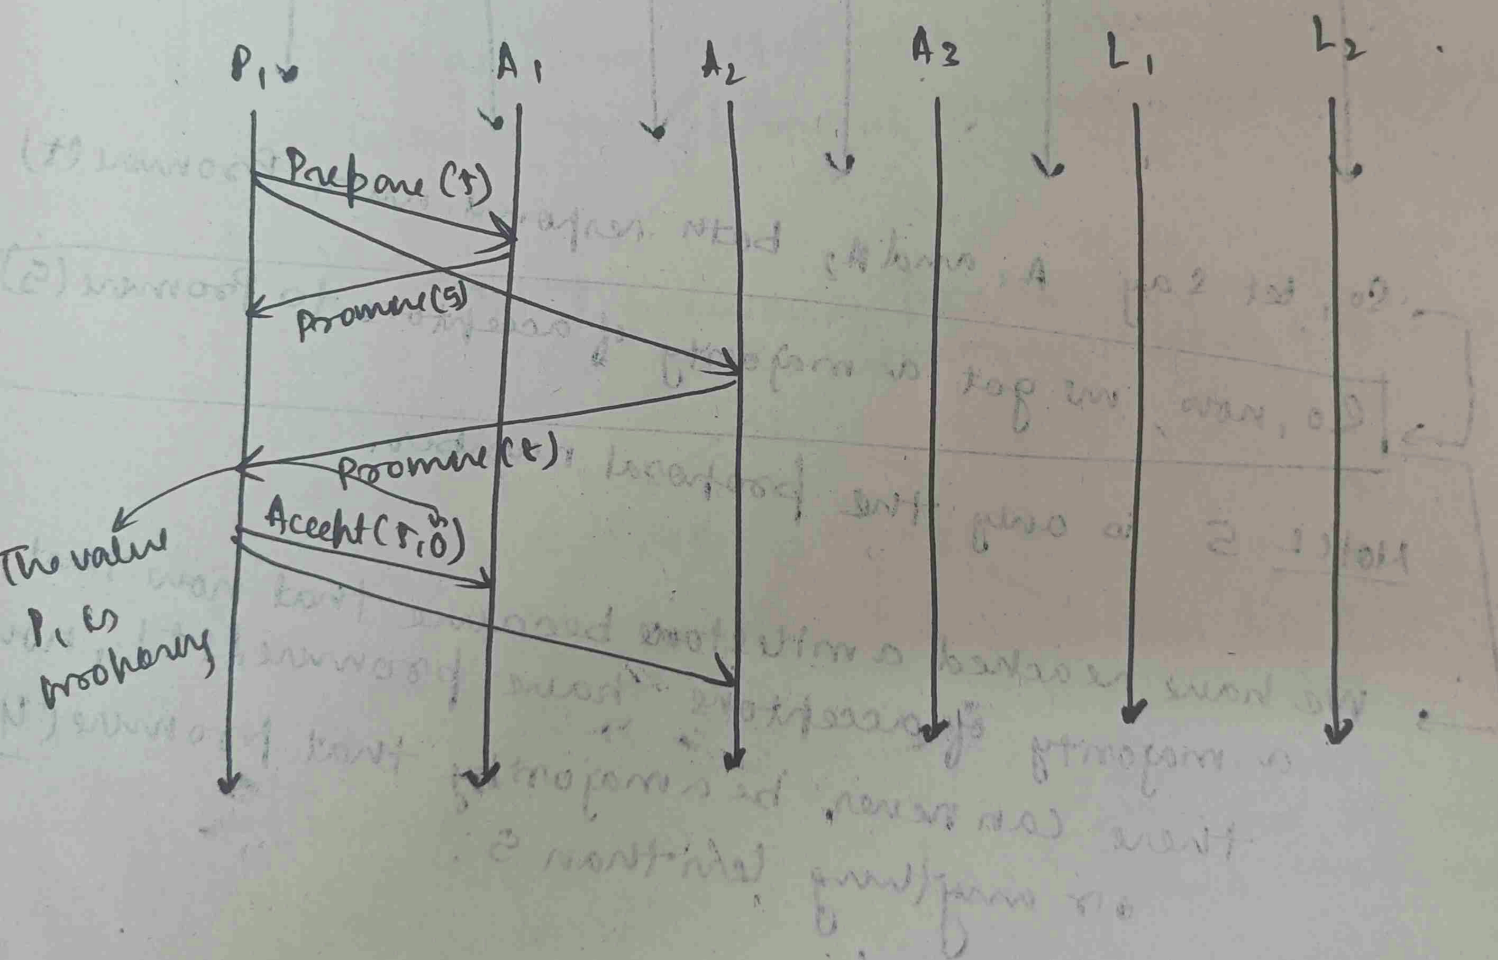
then it ignore it      It now promises  
 to ignore any  
 request that has a  
 proposal number  
 lower than that.

So, the Acceptors replies Proposer with a  
 "Promise" msg

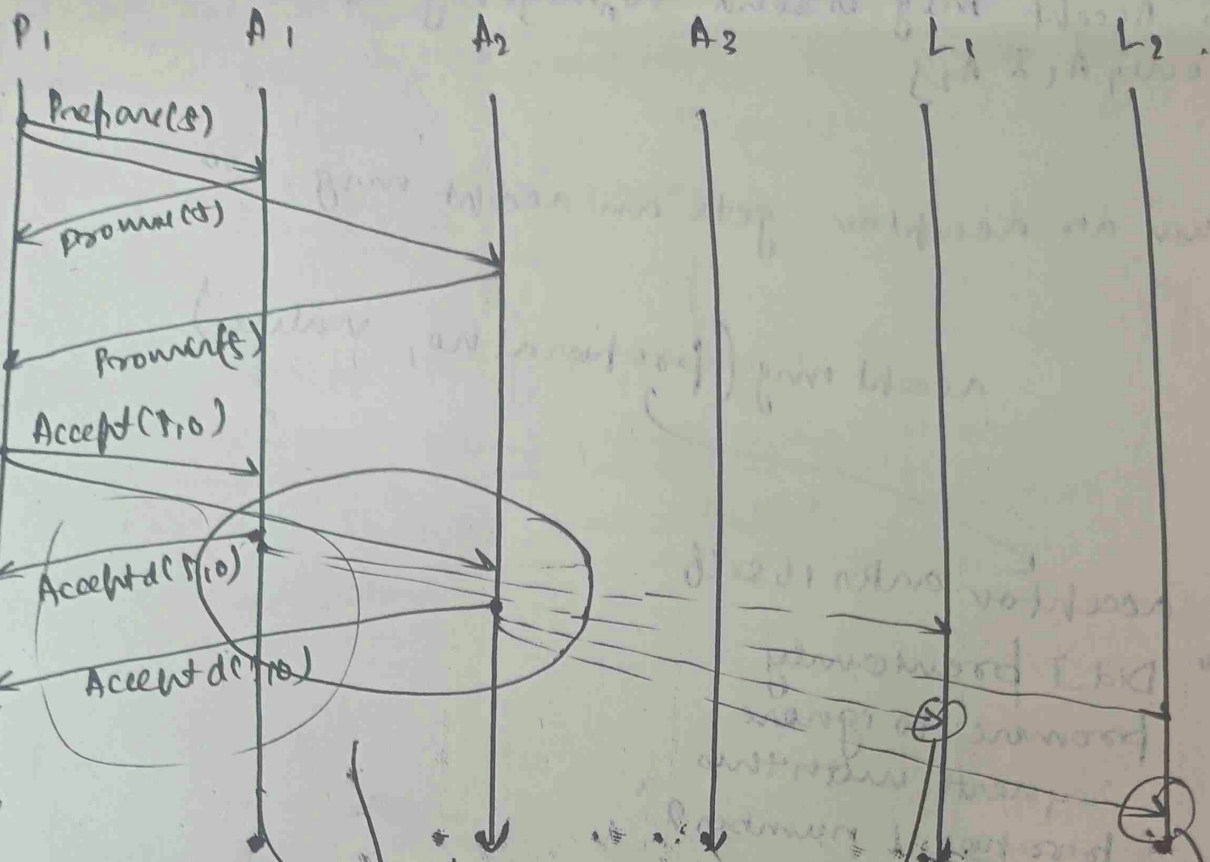
Promise(5)

Now that the proposer has received promise messages from a majority of acceptors, it can now send what's called an "accept msg" to a majority of acceptors.

Note 1: The Accept msg is going to have the 'proposal number that was promised' and the 'actual value that  $P_1$  wants to  $E$  propose'.







So, the point when a majority of acceptors send their {Accepted} msg for a particular proposal no. & particular value in the heart when consensus is reached.

Consensus is reached

Here P<sub>1</sub> knows that consensus is reached (when majority reply accepted)

L<sub>2</sub> knows here.

L<sub>1</sub> knows here

"So there is a difference b/w the moment at  
which consensus is reached and the  
moment at which everybody  
finds out that consensus is reached"